

Application of **Subset Selection** in **Efficient Machine Learning**

Maitreyee Research Showcase
Kiran Purohit


Advisor: **Prof. Sourangshu Bhattacharya**

Dept. of Computer Science & Engineering
IIT Kharagpur



Outline

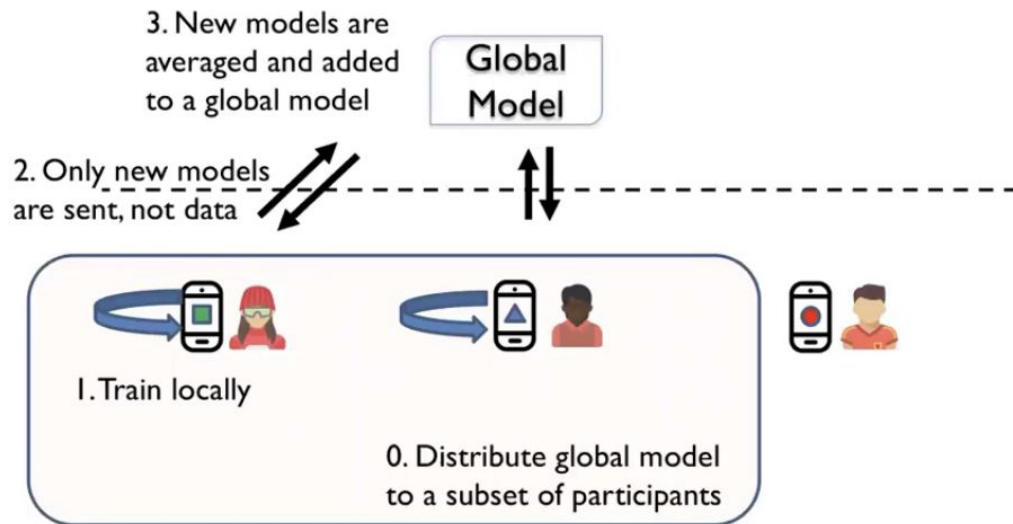
1. A Greedy Hierarchical Approach to Whole-Network Filter-Pruning in CNNs (*TMLR 2024*)
2. Efficient Exemplar Subset Selection for Complex Reasoning (*submitted*)



MISTRAL
AI_





Federated Learning

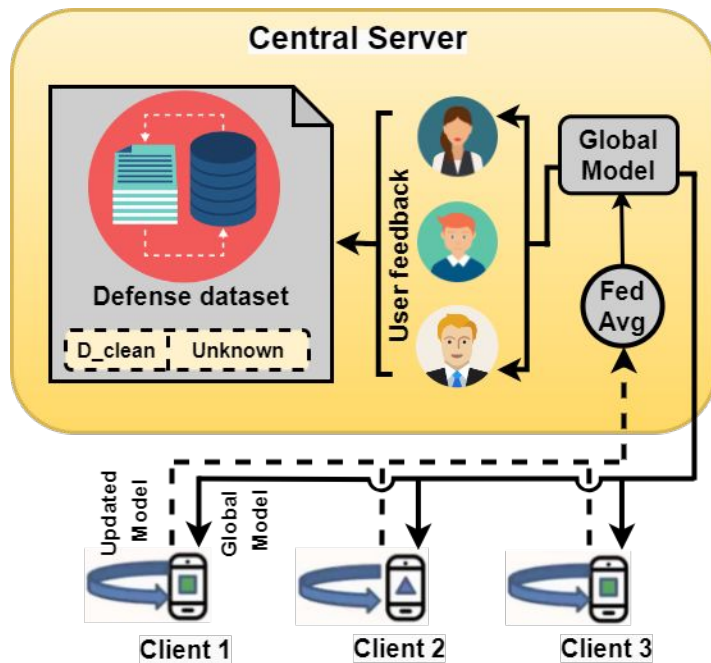


Backdoor Attacks

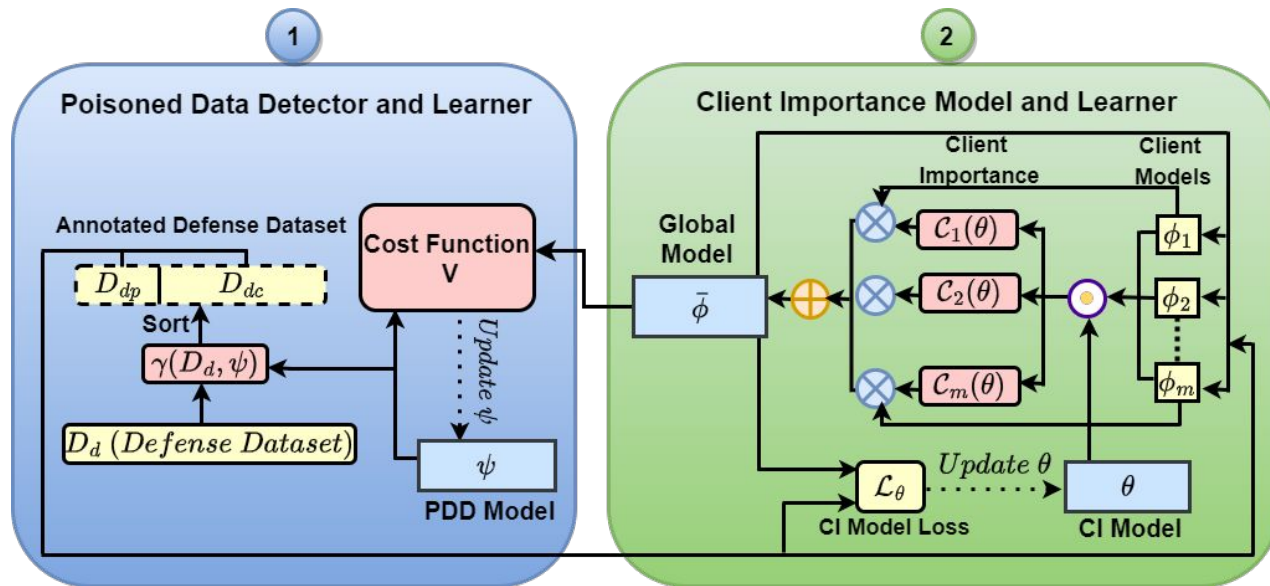
- ▶ Subtype of data poisoning
- ▶ Images with certain features are labeled differently
- ▶ Backdoor features can be artificial or natural
- ▶ Overall classification accuracy remains the same



DataDefense



Architecture of DataDefense



Filter Pruning

Burden of CNNs ——ResNet-152

60.2 million parameters and
231MB **storage** spaces;

380MB **memory** footprint


11.3 billion float point
operations (**FLOPs**).

Filter Pruning ——Benefits

reduces the **storage**
usage


decreases the **memory**
footprint

accelerates the inference




Network Pruning

Given a pre-trained network $\Phi(\cdot)$, the goal is to compress the network while maintaining the high performance as much as possible by removing the unnecessary parameters.




Pre-trained original network $\Phi(\cdot)$




Final pruned network $\Phi'(\cdot)$

Filter Pruning



Uniform Pruning

- ❖ Prune filters uniformly from each layer
- ❖ Process each layer independently and sequentially.



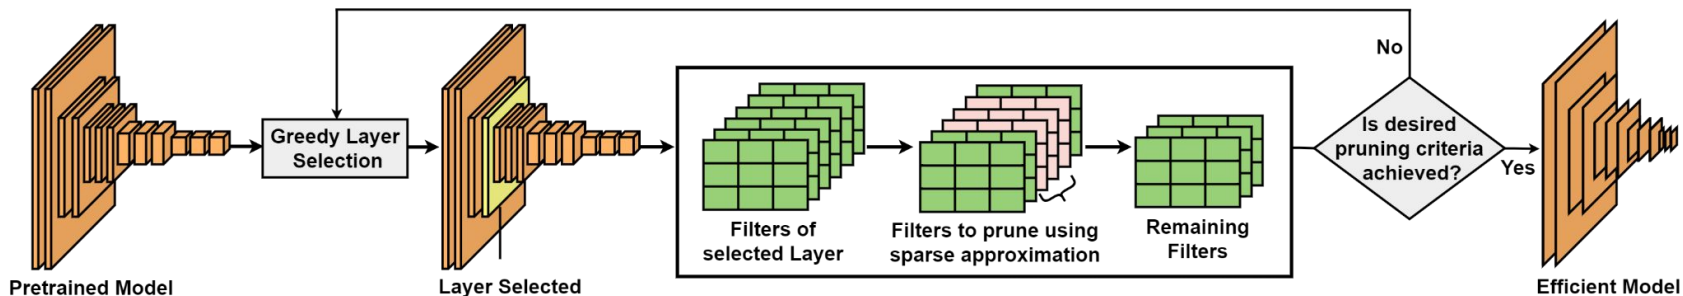
Non-Uniform Pruning

- ❖ Prune different fractions of filters from each layer
- ❖ All the layers in the network collectively make the final prediction

A Greedy Hierarchical Approach

- We developed faster **non-uniform pruning** methods.
- We used a hierarchical scheme with two-levels:
 - **filter pruning** - this step identifies the most appropriate filters to be pruned from each layer.
 - **layer selection** - this step selects the best layer to currently prune from.

We apply these two steps iteratively to achieve a non-uniform pruning.



Results and Analysis

Method	Test Acc (%)	Acc ↓ (%)	Param ↓ (%)	FLOPs ↓ (%)	VRAM (GB)
Dense RN16	92.1	0	-	-	7.62
Dense RN8	91.8	0	-	-	3.91
FP-Backward	92.9	-0.8	98.5	89.9	1.59
HBGS-B	93.0	-0.9	98.7	92.1	1.55
HBGTS-B	93.2	-1.1	98.8	94.3	1.51

Table: Comparison of pruning methods for ResNext101 32x16d (RN16) and a similar sized dense ResNext101 32x8d (RN8) on CIFAR10 at 98% parameter reduction.

- Our greedy hierarchical methods can be used for effectively pruning large models that exceed the capacity of commodity GPUs.
- ResNext101 32x16d has 193 M parameters and requires 7.62 GB of GPU memory for loading.
- We can efficiently deploy the pruned model on edge devices with GPU memory less than 2GB.

In-Context Learning (ICL)




Fig: Block Diagram of ICL

Exemplars / In-context examples / demonstration samples

<Question, Explanation, Answer>

Prompt

Instruction: You are a helpful, respectful and honest assistant helping to solve math word problems or tasks requiring reasoning or math. Follow given examples and solve the problems in step by step manner.

Exemplars :

[Question]: *The average age of three boys is 45 years and their ages are in proportion 3:5:7. What is the age in years of the youngest boy?*

[Explanation]: $3x + 5x + 7x = 45$,

$x = 3$,

$3x = 9$

[Answer]: **The answer is 9**

...


...

Test Input : Question:

Explanation: [INS] Answer: [INS]



Meta




MISTRAL
AI_

Explore-Exploit Paradigm

Exemplars

p_1 : While purchasing groceries ram bought **5 apples** ...

p_2 : Ephraim has **two machines** that make necklaces ...



Results and Analysis

Method	GSM8K	AquaRat	TabMWP	FinQA	StrategyQA
GPT-3.5-turbo					
dynamic					
KNN (Rubin et al., 2022)	53.45	51.96	77.07	51.52	81.83
KNN (S-BERT) (Rubin et al., 2022)	53.07	52.75	77.95	52.65	81.83
MMR (Ye et al., 2023b)	54.36	51.18	77.32	49.87	82.86
KNN+SC (Wang et al., 2023c)	80.21	62.59	83.08	54.49	83.88
MMR+SC (Wang et al., 2023c)	78.01	59.45	81.36	50.74	83.88
PromptPG (Lu et al., 2023b)	-	-	68.23	53.56	-
static					
Zero-Shot COT (Kojima et al., 2023)	67.02	49.60	57.10	47.51	59.75
Manual Few-Shot COT (Wei et al., 2023)	73.46	44.88	71.22	52.22	73.06
Random	67.79	49.80	55.89	53.70	81.02
PS+ (Wang et al., 2023b)	59.30	46.00	-	-	-
Auto-COT (Zhang et al., 2023b)	57.10	41.70	-	-	71.20
GraphCut (Iyer and Bilmes, 2013)	66.19	47.24	60.45	52.31	80.00
FacilityLocation (Iyer and Bilmes, 2013)	68.61	48.43	67.66	36.79	81.63
LENS (Li and Qiu, 2023)	69.37	48.82	77.27	54.75	79.79
LENS+SC (Li and Qiu, 2023)	79.37	57.87	80.68	60.06	82.24
Our Approach					
EXPLORA	77.86(▲12.24%) †	53.54(▲9.67%) †	83.07(▲7.51%) †	59.46(▲8.60%) †	85.71(▲5.63%) †
EXPLORA+SC	86.35(▲24.48%) ‡	63.39(▲29.84%) ‡	85.52(▲10.68%) ‡	64.52(▲17.84%) ‡	87.14 (▲9.21%) †
EXPLORA+KNN+SC	85.14 (▲22.73%) ‡	62.20(▲27.41%) ‡	86.29(▲12.39%) ‡	65.12(▲18.94%) ‡	88.37(▲10.75%) †
EXPLORA+MMR+SC	86.13(▲24.16%) ‡	63.78(▲30.64%) ‡	86.96(▲12.54%) ‡	64.60(▲17.99%) ‡	87.55(▲9.73%) †
GPT-4o					
LENS (Li and Qiu, 2023)	76.19	64.56	86.34	69.31	92.85
EXPLORA	93.63	69.29	90.12	72.71	95.10

THANK YOU
FOR
YOUR ATTENTION!!!



<https://github.com/kiranpurohit/>



@kiranpurohit08